

Homework – 3

Fall 2024

Spark Streaming

Sahithya Arveti Nagaraju (Person number: 50559752, UBITName: sarvetin)

Implement k-mer count program in Spark Streaming

- Create a local StreamingContext with two execution threads and a batch interval of 10 seconds.

Answer:

Below is the code for creating StreamingContext with two execution threads and batch interval of 10 seconds:

```
# Creating Spark context with 2 execution threads and naming App as "KMerCountApp"
conf=SparkConf().setAppName("KMerCountApp").setMaster("local[2]")
sc=SparkContext(conf=conf)
# Setting 10 second batch interval
ssc=StreamingContext(sc,10)
```

- SparkConf().setAppName("KMerCountApp") sets the name of the Spark application as 'KMerCountApp'.
- .setMaster("local[2]") sets the Spark context to run locally with 2 execution threads.
- StreamingContext(sc,10) initializes a StreamingContext with a batch interval of 10 seconds, which indicates that the system will handle and process the incoming data in chunks of 10 seconds.

DAG Visualization:



The above DAG shows a Spark Streaming job that continuously reads text data from a socket stream, processes it in batches, and writes the results to a destination. The windowing operation allows for real-time processing of the data stream, and the partitioning and mapping operations enable efficient parallel processing.

9	Streaming job from [output operation 0, batch time 18:28:50] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:52	16 ms	1/1	1/1
8	Streaming job from [output operation 0, batch time 18:28:50] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:51	0.6 s	1/1 (1 skipped)	1/1 (1 skipped)
7	Streaming job from [output operation 0, batch time 18:28:50] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:50	2 s	2/2	2/2
6	Streaming job from [output operation 0, batch time 18:28:40] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:42	20 ms	1/1	1/1
5	Streaming job from [output operation 0, batch time 18:28:40] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:41	0.7 s	1/1 (1 skipped)	1/1 (1 skipped)
4	Streaming job from [output operation 0, batch time 18:28:40] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:40	2 s	2/2	2/2
3	Streaming job from [output operation 0, batch time 18:28:30] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:33	19 ms	1/1	1/1
2	Streaming job from [output operation 0, batch time 18:28:30] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:32	0.6 s	1/1 (1 skipped)	1/1 (1 skipped)
1	Streaming job from [output operation 0, batch time 18:28:30] call at C:\Users\siria\AppData\Local\Programs\Python\Python39\lib\site-packages\py4j\clientserver.py:617	2024/12/01 18:28:30	2 s	2/2	2/2

Above Spark UI shows that each batch is processed every 10 seconds.

• **Create a DStream that represents streaming data from a TCP source (localhost:9999) [hint: use Netcat]**

Answer:

The DStream is created by connecting to a TCP source on localhost on port 9999. This is done by the following line in the code:

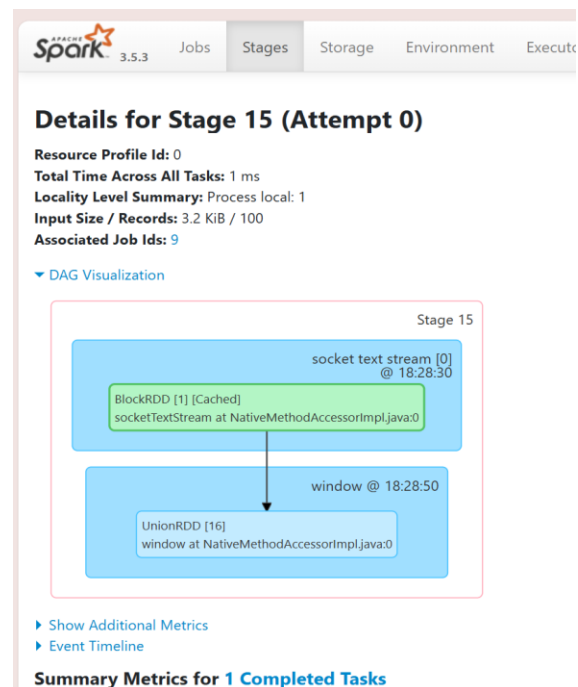
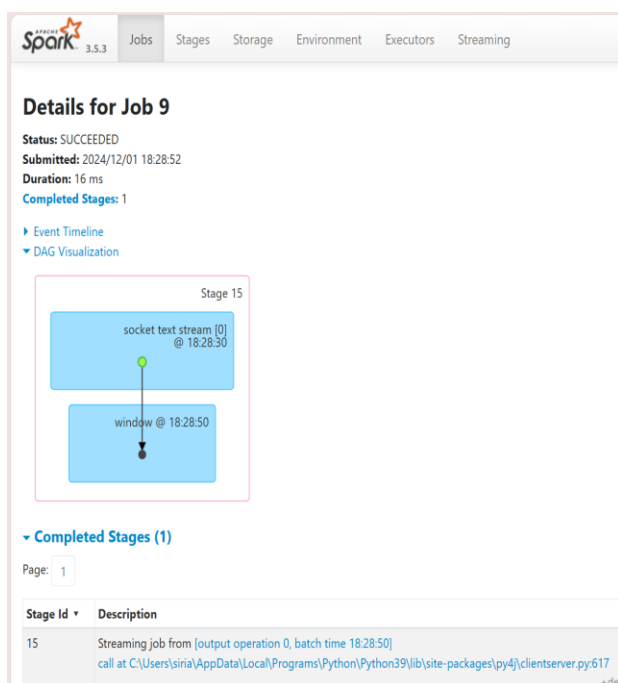
```
# DStream for receiving data from localhost (port 9999)
text_lines=ssc.socketTextStream("localhost",9999)
```

ssc.socketTextStream("localhost", 9999) sets up a DStream that listens for text data from a TCP socket on localhost at port 9999. The data from sentences.txt is streamed into this port using Netcat.

Netcat command for data streaming of sentences.txt into port 9999:

```
C:\Users\siria\DIC_Homework3>ncat -lk 9999 <sentences.txt
```

DAG Visualization:



Above DAG Visualization shows that a socket text stream (Dstream) created to stream data from TCP source (localhost:9999) with a window of 30 seconds with a sliding interval of 10 seconds.

- **Generate k-mers of length 3 from each line of text.**

Answer:

The create_kmers function in the code is responsible for generating k-mers of length 3 from each line of text:

```

5 # Creating k-mers of Length 3
6 def create_kmers(text):
7     k = 3
8     return list(map(lambda i:text[i:i+k], range(len(text)- k+1)))
9

```

The function `create_kmers` takes a string `text` as input and generates a list of substrings (k-mers) of length 3. It uses the `map` function to iterate through all possible starting indices `i` of length-3 substrings within the text. The range `range(len(text) - k + 1)` ensures that the function doesn't attempt to create a k-mer starting too close to the end of the string. It generates a list of all the k-mers of length 3 which are possible in the text provided as shown in below screenshot:

The screenshot shows a Jupyter Notebook interface with the file `output_kmers.txt` and a last checkpoint 4 hours ago. The output displays 18 lines, each representing a text input and its corresponding list of 3-letter k-mers. For example, the first line is: "K-mers for line 'mqnslvaykpmwtficftyknjbwdzkl': mqn qns nsl slv lvv vva vay ayk ykp kpm pmw mwt wtv tvf vfi fic icf cft fty tyk ykn knj njb jbw bwd wdz dzk zkl". The output continues with similar lines for various other text inputs, showing the first few k-mers for each.

Also, saved all the k-mers for each line in a text file named `output_kmers.txt` for reference.

- **Count the occurrences of each k-mer and print the k-mer counts. Take screenshots of your output.**

Answer:

The k-mer frequencies are counted by using the `reduceByKey` operation in the `process_kmers` function:

```

15  # Generating k-mers of length 3 for each line in the RDD and count k-mers
16  kmer_list = rdd.flatMap(lambda line:create_kmers(line))
17  kmer_count= kmer_list.map(lambda kmer:(kmer,1)).reduceByKey(lambda x,y:x+y)

```

- `rdd.flatMap(lambda line: create_kmers(line))`: Transforms every line of the RDD to a list of k-mers. The `flatMap` function will flatten all the lists of k-mers into one single RDD.
- `reduceByKey(lambda x, y: x + y)`: reduces the RDD by key - i.e. k-mer - and sums up values, so counts a frequency of every k-mer.

The first top 10 k-mer counts are printed using the following code:

```

26  # Sorting the k-mers by frequency and taking the top 10 k-mers
27  top_kmers = kmer_count.takeOrdered(10, key=lambda x: -x[1])
28
29  print("\nTop 10 K-mer counts:")
30  for kmer, freq in top_kmers:
31      print(f"Count of '{kmer}': {freq}")

```

Explanation:

The `kmer_count.collect()` method collects all the k-mer counts as tuples of the form (kmer, frequency). A for loop iterates over these tuples, printing the first 10 k-mers and their frequencies.

```
Processing batch at: 2024-12-01 23:09:20
```

```
Top 10 K-mer counts:
```

```
Count of 'eha': 4  
Count of 'xhn': 4  
Count of 'byb': 3  
Count of 'kqx': 3  
Count of 'cgy': 3  
Count of 'mkl': 3  
Count of 'lpb': 3  
Count of 'kfn': 3  
Count of 'bhk': 3  
Count of 'ykp': 2
```

```
Processing batch at: 2024-12-01 23:09:30
```

```
Top 10 K-mer counts:
```

```
Count of 'eha': 4  
Count of 'xhn': 4  
Count of 'byb': 3  
Count of 'kqx': 3  
Count of 'cgy': 3  
Count of 'mkl': 3  
Count of 'lpb': 3  
Count of 'kfn': 3  
Count of 'bhk': 3  
Count of 'ykp': 2
```

```
Processing batch at: 2024-12-01 23:09:40
```

```
Top 10 K-mer counts:
```

```
Count of 'eha': 4  
Count of 'xhn': 4  
Count of 'byb': 3  
Count of 'kqx': 3  
Count of 'cgy': 3  
Count of 'mkl': 3  
Count of 'lpb': 3  
Count of 'kfn': 3  
Count of 'bhk': 3  
Count of 'ykp': 2
```

The above is the output of kmer counts, here I'm only displaying only top 10 kmer counts since there are a lot of kmers. Additionally, it can be observed that Dstream is performing kmer counting for every 10 seconds interval.

Also, saved all the k-mers count in a text file named output_kmer_count.txt for reference.

```

1 mqn: 1
2 qns: 1
3 nsl: 1
4 slv: 1
5 lvv: 1
6 vva: 1
7 vay: 1
8 ayk: 1
9 ykp: 2
10 kpm: 1
11 pmw: 2
12 mwt: 1
13 wtv: 1
14 tvf: 1
15 vfi: 1
16 fic: 1
17 icf: 1
18 cft: 1
19 fty: 1
20 tyk: 2
21 ykn: 1
22 knj: 1
23 njb: 1
24 jbw: 1
25 bwd: 1
26 wdz: 1
27 dzk: 1
28 zkl: 1
29 nkp: 1
30 kpz: 2
31 pzo: 1
32 zov: 1
33 ova: 1
34 vac: 1
35 acr: 2
36 crm: 1
37 rmb: 1

```

Alternative method:

The above method using netcat was getting whole data from sentences.txt at first time interval itself, so it was performing kmer operation at first time interval and was not getting any data streamed after that from netcat.


For alternative method, first run tcp.py file in one terminal like below screenshot and then run code.py in another terminal

```

PS C:\Users\siria\DIC_Homework3> python .\tcp.py
Server is listening on port 9999...
Connection from ('127.0.0.1', 35308)
Connection from ('127.0.0.1', 35711)
Connection from ('127.0.0.1', 36107)
Connection from ('127.0.0.1', 36514)
Connection from ('127.0.0.1', 36950)
Connection from ('127.0.0.1', 37378)
Connection from ('127.0.0.1', 37757)
Connection from ('127.0.0.1', 38140)
Connection from ('127.0.0.1', 38542)
Connection from ('127.0.0.1', 38936)
Connection from ('127.0.0.1', 39343)
Connection from ('127.0.0.1', 39722)
Connection from ('127.0.0.1', 40085)
Connection from ('127.0.0.1', 40454)

```

So, I have used an alternative method for sending data using TCP socket program as shown below to stream data in localhost 9999,


Jupyter

tcp.py
Last Checkpoint: 51 minutes ago

File
Edit
View
Settings
Help

```

1 import socket
2 import time
3
4 # Create a TCP/IP socket
5 server_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
6 server_socket.bind(('localhost', 9999))
7 server_socket.listen(1)
8
9 print("Server is listening on port 9999...")
10
11 while True:
12     connection, client_address = server_socket.accept()
13     try:
14         print(f"Connection from {client_address}")
15         # Continuously send lines of data
16         with open('sentences.txt', 'r') as file:
17             for line in file:
18                 connection.sendall(line.encode())
19                 time.sleep(1) # Simulate streaming with a delay
20     finally:
21         connection.close()
22

```

By using above method, data was sending continuously and able to get the kmer counts in batches. Here are the few screenshots of kmer counts after using this method,


```
Processing batch at: 2024-12-01 22:42:59
24/12/01 22:42:51 WARN RandomBlockReplicationPolicy: Expecting 1 replicas with only 0 peer/s.
24/12/01 22:42:51 WARN BlockManager: Block input-0-1733110971000 replicated to only 0 peer(s) instead of 1 peers
24/12/01 22:42:52 WARN RandomBlockReplicationPolicy: Expecting 1 replicas with only 0 peer/s.
24/12/01 22:42:52 WARN BlockManager: Block input-0-1733110972000 replicated to only 0 peer(s) instead of 1 peers
24/12/01 22:42:53 WARN RandomBlockReplicationPolicy: Expecting 1 replicas with only 0 peer/s.
24/12/01 22:42:53 WARN BlockManager: Block input-0-1733110973000 replicated to only 0 peer(s) instead of 1 peers
24/12/01 22:42:54 WARN RandomBlockReplicationPolicy: Expecting 1 replicas with only 0 peer/s.
24/12/01 22:42:54 WARN BlockManager: Block input-0-1733110974000 replicated to only 0 peer(s) instead of 1 peers
24/12/01 22:42:55 WARN RandomBlockReplicationPolicy: Expecting 1 replicas with only 0 peer/s.
24/12/01 22:42:55 WARN BlockManager: Block input-0-1733110975000 replicated to only 0 peer(s) instead of 1 peers
24/12/01 22:42:56 WARN RandomBlockReplicationPolicy: Expecting 1 replicas with only 0 peer/s.
24/12/01 22:42:56 WARN BlockManager: Block input-0-1733110976000 replicated to only 0 peer(s) instead of 1 peers

Top 10 K-mer counts:
Count of 'lvv': 1
Count of 'ayk': 1
Count of 'tvf': 1
Count of 'zkl': 1
Count of 'nkp': 1
Count of 'rmb': 1
Count of 'bit': 1
Count of 'vkp': 1
Count of 'yjb': 1
Count of 'zly': 1
```

Processing batch at: 2024-12-01 22:43:00

```
Top 10 K-mer counts:
Count of 'cst': 2
Count of 'irw': 2
Count of 'ykp': 2
Count of 'kpx': 2
Count of 'ayk': 1
Count of 'tvf': 1
Count of 'nkp': 1
Count of 'rmb': 1
Count of 'bit': 1
Count of 'yjb': 1
```

Processing batch at: 2024-12-01 22:50:51

```
Top 10 K-mer counts:
Count of 'xhp': 2
Count of 'bhk': 2
Count of 'sjq': 2
Count of 'amw': 2
Count of 'scs': 2
Count of 'lqy': 2
Count of 'jcd': 2
Count of 'hno': 2
Count of 'dlq': 2
Count of 'amk': 2
```

Processing batch at: 2024-12-01 22:51:41

```
Top 10 K-mer counts:
Count of 'kpx': 2
Count of 'lqy': 2
Count of 'ogw': 2
Count of 'ykp': 2
Count of 'kna': 2
Count of 'cxb': 2
Count of 'hno': 2
Count of 'znh': 2
Count of 'scs': 2
Count of 'kqx': 2
```

Processing batch at: 2024-12-01 22:52:32

Top 10 K-mer counts:

Count of 'cst': 2

Count of 'irw': 2

Count of 'yjh': 2

Count of 'byb': 2

Count of 'bvn': 2

Count of 'nat': 2

Count of 'njj': 2

Count of 'jbx': 2

Count of 'cxb': 2

Count of 'yab': 2

2024-12-01 22:53:07 WAPN-Random

Processing batch at: 2024-12-01 22:53:25

Top 10 K-mer counts:

Count of 'yda': 2

Count of 'fyd': 2

Count of 'kzd': 2

Count of 'amj': 2

Count of 'bvn': 2

Count of 'nat': 2

Count of 'cjj': 2

Count of 'njj': 2

Count of 'jbx': 2

Count of 'tzs': 2

2024-12-01 22:54:02 WAPN-Random

Processing batch at: 2024-12-01 22:54:20

Top 10 K-mer counts:

Count of 'wwf': 2

Count of 'byb': 2

Count of 'amj': 2

Count of 'cjj': 2

Count of 'mkl': 2

Count of 'tzs': 2

Count of 'qwh': 2

Count of 'qcl': 2

Count of 'tlr': 2

Count of 'wyu': 2

2024-12-01 22:54:56 WAPN-Random